

FUNCTIONAL PROGRAMMING

Gerko Vink 

g.vink@uu.nl

Methodology & Statistics @ Utrecht University

9 Jun 2025

DISCLAIMER

I owe a debt of gratitude to many people as the thoughts and code in these slides are the process of years-long development cycles and discussions with my team, friends, colleagues and peers. When someone has contributed to the content of the slides, I have credited their authorship.

Images are either directly linked, or generated with StableDiffusion or DALL-E. That said, there is no information in this presentation that exceeds legal use of copyright materials in academic settings, or that should not be part of the public domain.

Warning

You **may use** any and all content in this presentation - including my name - and submit it as input to generative AI tools, with the following **exception**:

- You must ensure that the content is not used for further training of the model

SLIDE MATERIALS AND SOURCE CODE

Materials

- lecture slides on Moodle
- course page: www.gerkovink.com/sur
- source: github.com/gerkovink/sur

RECAP

Gisteren hebben we deze onderwerpen behandeld:

- Beschrijvende statistiek
- Kruistabellen en frequentieverdelingen
- χ^2 -toets
- Andere toets- en associatiematen
- Simpele lineaire regressie
- Analyses draaien op groepen

TODAY

Vandaag behandelen we de volgende onderwerpen:

- Zelf functies ontwikkelen, gebruiken en debuggen
- Map / Reduce workflows
- Binaire operators
- Trekken uit verdelingen
- Random number generation

PACKAGES WE USE

```
1 library(dplyr)      # data manipulation
2 library(purrr)      # functional programming
3 library(furrr)      # parallel processing
4 library(magrittr)   # flexible pipes
5 library(mice)       # for the boys data
6
7 # fix the random seed
8 set.seed(123)
```

WRITING FUNCTIONS

WRITE YOUR OWN FUNCTION

The function is `function()`

```
1 my_function <- function(arguments) {  
2  
3   expressions  
4  
5   return(output)  
6  
7 }
```

- `arguments` are input of the function
- `expressions` are operations performed on the arguments
- `output` an object containing the output (e.g. vector, matrix, list, etc.)
- `return` explicit return of the output (optional, but recommended!)

```
1 my_function <- function(arguments) {  
2  
3   expressions  
4  
5   output # less clear that this is returned  
6  
7 }
```


TOSSING A DIE

A function without arguments that simulates tossing a die

```
1 die <- function() {  
2   # throw die  
3   eyes <- sample(1:6, size = 1)  
4   # return the outcome  
5   return(eyes)  
6 }
```

```
1 c(die(), die(), die())
```

```
[1] 3 6 3
```

DEFINING AN ARGUMENT

The argument `n` specifies the number of throws of the die

```
1 dice <- function(n) {  
2   # n is the number of dice to toss  
3   # replace = TRUE allows for repeated outcomes  
4   # returns a vector of length n  
5   return(sample(1:6, size = n, replace = TRUE))  
6 }  
7  
8 dice(10)
```

```
[1] 2 2 6 3 5 4 6 6 1 2
```

MULTIPLE RETURNS

A function `dice(n)` returning a list with

- the outcomes of the `n` throws, their frequencies and their mean

```
1 dice <- function(n) {
2   # throw dice n times
3   eyes <- sample(1:6, size = n, replace = TRUE)
4   # prepare structured output
5   return(list(outcomes = eyes,
6               freqs     = table(eyes),
7               mean      = mean(eyes)))
8 }
9 dice(10)
```

\$outcomes

[1] 3 5 3 3 1 4 1 1 5 3

\$freqs

eyes

1 3 4 5

3 4 1 2

\$mean

[1] 2.9

DEFAULT ARGUMENTS

The default is a fair die (each outcome has probability 1/6)

- the user can change this if so desired

```
1 dice <- function(n, p = rep(1/6, 6)) {
2   # throw dice n times with probability p
3   eyes <- sample(1:6, size = n, replace = TRUE, prob = p)
4   # prepare structured output
5   return(list(outcomes = eyes,
6               frequency = table(eyes),
7               mean      = mean(eyes)))
8 }
9 dice(100)
```

\$outcomes

```
[1] 5 5 3 2 1 1 6 6 2 4 6 3 3 3 2 4 4 4 2 2 3 4 3 1 2 4 6 2 5 3 2 6 1 4 5 2 4
[38] 3 6 4 6 6 6 4 6 5 6 2 4 3 4 5 4 2 3 6 4 6 2 4 1 1 1 3 2 5 4 5 3 3 6 2 4 5
[75] 5 3 4 1 4 1 1 5 4 2 1 3 2 1 6 2 5 1 5 4 5 3 3 3 4 1
```

\$frequency

eyes

```
1  2  3  4  5  6
14 17 18 22 14 15
```

\$mean

```
[1] 3.5
```

UNFAIR DICE

The following command throws 100 unfair dice

- probabilities for rolling a 1, 2, 3, 4, 5 is 0.1
- probability for rolling a 6 is 0.5

```
1 c(rep(.1, 5), .5)
```

```
[1] 0.1 0.1 0.1 0.1 0.1 0.5
```

```
1 dice(100, p = c(rep(.1, 5), .5))
```

\$outcomes

```
[1] 6 6 6 4 4 2 4 5 3 4 2 5 1 6 6 6 6 6 2 6 6 6 6 5 2 6 6 6 6 6 3 6 6 6 3 6 4
[38] 6 6 3 5 6 6 6 4 6 2 5 4 4 6 3 2 3 2 6 5 6 3 6 6 3 1 1 6 6 1 4 1 6 6 4 6 3
[75] 6 1 4 3 6 2 6 6 6 6 6 6 6 4 6 5 6 6 2 1 6 1 5 4 6 6
```

\$frequency

eyes

```
1  2  3  4  5  6
8  9 10 13  8 52
```

\$mean

```
[1] 4.6
```

APPLYING YOUR FUNCTION

apply()

The `apply()` function is used to apply a function to the rows or columns of a matrix or array. It takes three main arguments: the data, the margin (1 for rows, 2 for columns), and the function to apply.

```
1 calc_mean <- function(x) {
2   return(mean(x, na.rm = TRUE))
3 }
4 # select random 10 rows from the numeric columns of boys
5 numboys <- boys %>% select(where(is.numeric))
6 which_rows <- sample(1:nrow(numboys), 10)
7 numboys <- numboys[which_rows, ]
8 # over the columns
9 apply(numboys, FUN = calc_mean, MARGIN = 2)
```

age	hgt	wgt	bmi	hc	tv
15.0746	168.3200	62.7300	21.2590	55.1300	10.5000

```
1 # over the rows
2 apply(numboys, FUN = calc_mean, MARGIN = 1)
```

5975	7062	6131	5897	4505	7073	7088	4487
58.03867	67.37580	73.66980	68.45680	52.04250	62.35783	71.46680	48.28217
6693	2423						
72.43360	35.15900						

lapply()

`lapply()` does the same as `apply()`, but it is used for lists. It applies a function to each element of a list and returns a list of results.

```
1 lapply(numboys, FUN = calc_mean)
```

\$age

[1] 15.0746

\$hgt

[1] 168.32

\$wgt

[1] 62.73

\$bmi

[1] 21.259

\$hc

[1] 55.13

\$tv

[1] 10.5

sapply()

`sapply()` does the same as `lapply()`, but it simplifies the output to a vector or matrix if possible. It is useful when you want to avoid dealing with lists.

```
1 sapply(numboys, FUN = calc_mean)
```

age	hgt	wgt	bmi	hc	tv
15.0746	168.3200	62.7300	21.2590	55.1300	10.5000

tapply()

`tapply()` is used to apply a function to subsets of a vector, grouped by one or more factors. It is particularly useful for summarizing data based on grouping variables.

```
1 tapply(boys$hgt, boys$reg, FUN = calc_mean)
```

```
   north    east    west    south    city
151.6316 133.9648 130.2783 128.0022 125.8577
```

```
1 boys %>%
2   group_by(reg) %>%
3   summarise(mean_hgt = mean(hgt, na.rm = TRUE))
```

```
# A tibble: 6 × 2
  reg    mean_hgt
<fct>    <dbl>
1 north    152.
2 east    134.
3 west    130.
4 south    128.
5 city    126.
6 <NA>     73.0
```

MAP / REDUCE

map()

The `map()` function is part of the `purrr` package, which is designed for functional programming in R. It allows you to apply a function to each element of a list or vector, returning a list of results.

```
1 boys %>%
2   split(.$reg) %>% # split the data by region
3   map(~ lm(hgt ~ age, data = .x) %>% # map the linear model function
4     coef()) # extract coefficients
```

```
$north
(Intercept)      age
  74.104664    6.376882
```

```
$east
(Intercept)      age
  73.229714    6.507535
```

```
$west
(Intercept)      age
  69.446550    6.646496
```

```
$south
(Intercept)      age
  70.410123    6.566541
```

```
$city
(Intercept)      age
  69.010565    6.724608
```

split()

```
1 boys %>%
2   split(.$reg)
```

\$north

	age	hgt	wgt	bmi	hc	gen	phb	tv	reg
127	0.093	56.0	5.410	17.25	40.0	<NA>	<NA>	NA	north
198	0.117	57.0	5.260	16.18	40.0	<NA>	<NA>	NA	north
238	0.142	58.0	5.220	15.51	40.1	<NA>	<NA>	NA	north
248	0.147	57.3	4.950	15.07	36.8	<NA>	<NA>	NA	north
873	0.594	70.8	8.970	17.89	45.2	<NA>	<NA>	NA	north
911	0.673	71.0	9.000	17.85	46.5	<NA>	<NA>	NA	north
1212	0.996	77.1	10.390	17.47	47.1	<NA>	<NA>	NA	north
1278	1.040	77.5	9.300	15.48	46.3	<NA>	<NA>	NA	north
1511	1.292	79.0	10.700	17.14	47.3	<NA>	<NA>	NA	north
1617	1.481	NA	12.040	NA	47.5	<NA>	<NA>	NA	north
1684	1.530	80.0	10.785	16.85	46.1	<NA>	<NA>	NA	north
1877	1.793	86.0	13.400	18.11	47.0	<NA>	<NA>	NA	north
1882	1.798	81.8	10.535	15.74	46.7	<NA>	<NA>	NA	north
1927	1.848	NA	13.200	NA	50.0	<NA>	<NA>	NA	north
2044	2.020	88.3	13.000	16.67	50.0	<NA>	<NA>	NA	north
2168	2.198	94.2	14.980	16.88	51.0	<NA>	<NA>	NA	north
2313	2.576	86.0	10.700	14.46	49.5	<NA>	<NA>	NA	north

map()

The `map()` function is particularly useful for iterating over lists or vectors and applying a function to each element. It can be used to perform operations like calculations, transformations, or data extraction.

```
1 out <- boys %>%
2   split(.$reg) %>% # split the data by region
3   map(~ lm(hgt ~ age, data = .x) %>% # map the linear model function
4       coef()) # extract coefficients
5
6 is.list(out)
```

```
[1] TRUE
```

```
1 names(out)
```

```
[1] "north" "east"  "west"  "south" "city"
```

```
1 out$city
```

```
(Intercept)      age
 69.010565    6.724608
```

map() ON LARGE LISTS

In the below example, we take a bootstrap sample (with replacement) from the `boys` data 1000 times, and then run a simple linear model on all 1000 bootstrap samples separately.

```
1 sample_rows <- function(x) {  
2   out <- x[sample(1:nrow(x), replace = TRUE), ]  
3   return(out)  
4 }  
5 samples <- replicate(n = 1000, expr = sample_rows(boys), simplify = FALSE)  
6 samples_lm <-  
7   samples %>%  
8   map(~.x %$%  
9     lm(hgt ~ age) %>%  
10    coef())
```

map() ON LARGE LISTS

```
1 # how many samples?
2 length(samples)
```

```
[1] 1000
```

```
1 # what is the first sample?
2 samples[[1]] %>% slice_head()
```

```
      age  hgt  wgt  bmi hc  gen  phb tv  reg
4501 12.342 158.1 54.9 21.96 NA <NA> <NA> NA west
```

```
1 # what is the first sample's linear model?
2 samples_lm[[1]]
```

```
(Intercept)      age
 70.975166    6.573117
```

```
1 # what are the first three samples' linear model?
2 samples_lm[1:3]
```

```
[[1]]
(Intercept)      age
 70.975166    6.573117
```

```
[[2]]
(Intercept)      age
 69.902756    6.687102
```

```
[[3]]
(Intercept)      age
 70.595459    6.583677
```


Reduce()

With `reduce()`, you can combine the results obtained with `map()`.

```
1 reduce(samples_lm, `+`) # sum of lm coefficients
```

```
(Intercept)      age  
 70681.662    6594.446
```

```
1 reduce(samples_lm, `+`) / 1000 # average lm coefficients
```

```
(Intercept)      age  
 70.681662    6.594446
```

STRUCTURING THE OUTPUT OF `map()`

`map()` returns a list, which can be structured into a data frame using `map_df()`. This is useful when you want to convert the results of `map()` into a tidy format.

```
1 # with map_df() instead of map() to return a data frame
2 samples %>%
3   map_df(~.x %$%
4           lm(hgt ~ age) %>%
5           coef())
```

```
# A tibble: 1,000 × 2
#   `(Intercept)` age
#   <dbl> <dbl>
1      71.0  6.57
2      69.9  6.69
3      70.6  6.58
4      70.8  6.56
5      70.2  6.56
6      70.2  6.67
7      70.8  6.67
8      72.2  6.48
9      69.7  6.71
10     70.9  6.57
# i 990 more rows
```

STRUCTURING THE OUTPUT OF `map()`

```
1 # with map_df() if an object already exists as a list
2 samples_lm %>%
3   map_df(~tibble(intercept = .x[1], slope = .x[2]))
```

A tibble: 1,000 × 2

	intercept	slope
	<dbl>	<dbl>
1	71.0	6.57
2	69.9	6.69
3	70.6	6.58
4	70.8	6.56
5	70.2	6.56
6	70.2	6.67
7	70.8	6.67
8	72.2	6.48
9	69.7	6.71
10	70.9	6.57

i 990 more rows

FUTURES

- The **future** package enables **asynchronous** and **parallel** processing in R.
- It allows R to perform tasks **in the background**, freeing up your current R session.
- Ideal for:
 - Speeding up long-running computations
 - Running tasks concurrently

WHY USE **future**?

- Normally, R runs code **line-by-line** (sequentially).
- **future** lets you run tasks **in parallel**, improving efficiency.
- Example use cases:
 - Simulations
 - Data processing across multiple cores
 - Web scraping multiple pages

future_map()

The `future_map()` function is part of the `furrr` package, which integrates the `future` package with the `purrr` package's mapping functions.

It allows you to apply a function to each element of a list or vector in parallel, making it easier to handle large datasets or computationally intensive tasks.

```
1 # Set up parallel processing
2 plan(multisession) # Use multiple cores for parallel processing
3 samples %>%
4   future_map_dfr(~.x %$%
5     lm(hgt ~ age) %>%
6     coef())
```

```
# A tibble: 1,000 × 2
#   `(Intercept)` age
#   <dbl> <dbl>
1      71.0  6.57
2      69.9  6.69
3      70.6  6.58
4      70.8  6.56
5      70.2  6.56
6      70.2  6.67
7      70.8  6.67
8      72.2  6.48
9      69.7  6.71
10     70.9  6.57
# i 990 more rows
```

```
1 plan(sequential) # Stop parallel processing and reset to sequential processing
```

future_map()

The `future_map()` function is part of the `furrr` package, which integrates the `future` package with the `purrr` package's mapping functions.

It allows you to apply a function to each element of a list or vector in parallel, making it easier to handle large datasets or computationally intensive tasks.

```
1 # Set up parallel processing
2 plan(multisession) # Use multiple cores for parallel processing
3 samples %>%
4   future_map_dfc(~.x %$%
5     lm(hgt ~ age) %>%
6     coef())
```

```
# A tibble: 2 × 1,000
  ...1 ...2 ...3 ...4 ...5 ...6 ...7 ...8 ...9 ...10 ...11 ...12 ...13
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 71.0  69.9  70.6  70.8  70.2  70.2  70.8  72.2  69.7  70.9  70.3  70.5  69.5
2  6.57  6.69  6.58  6.56  6.56  6.67  6.67  6.48  6.71  6.57  6.61  6.55  6.68
# i 987 more variables: ...14 <dbl>, ...15 <dbl>, ...16 <dbl>, ...17 <dbl>,
#   ...18 <dbl>, ...19 <dbl>, ...20 <dbl>, ...21 <dbl>, ...22 <dbl>,
#   ...23 <dbl>, ...24 <dbl>, ...25 <dbl>, ...26 <dbl>, ...27 <dbl>,
#   ...28 <dbl>, ...29 <dbl>, ...30 <dbl>, ...31 <dbl>, ...32 <dbl>,
#   ...33 <dbl>, ...34 <dbl>, ...35 <dbl>, ...36 <dbl>, ...37 <dbl>,
#   ...38 <dbl>, ...39 <dbl>, ...40 <dbl>, ...41 <dbl>, ...42 <dbl>,
#   ...43 <dbl>, ...44 <dbl>, ...45 <dbl>, ...46 <dbl>, ...47 <dbl>, ...
```

```
1 plan(sequential) # Stop parallel processing and reset to sequential processing
```

BINAIRY OPERATORS LIKE `%in%` AND `%>%`

HOW BINAIRY OPERATORS WORK

Binary operators are functions that take two arguments and return a single value. In R, you can create your own binary operators using the "%operator%" syntax.

```
1 `my_operator` <- function(x, y) {
2   # perform some operation on x and y
3   result <- x + y # example operation: addition
4   return(result)
5 }
6
7 4 my_operator 5
```

```
[1] 9
```

```
1 1:4 my_operator 5:8
```


```
[1] 6 8 10 12
```

Binary operators allow you to write the function in a more natural way, similar to mathematical notation. You can use them for various operations, such as addition, subtraction, multiplication, or any custom operation you define.

DRAWING FROM DISTRIBUTIONS

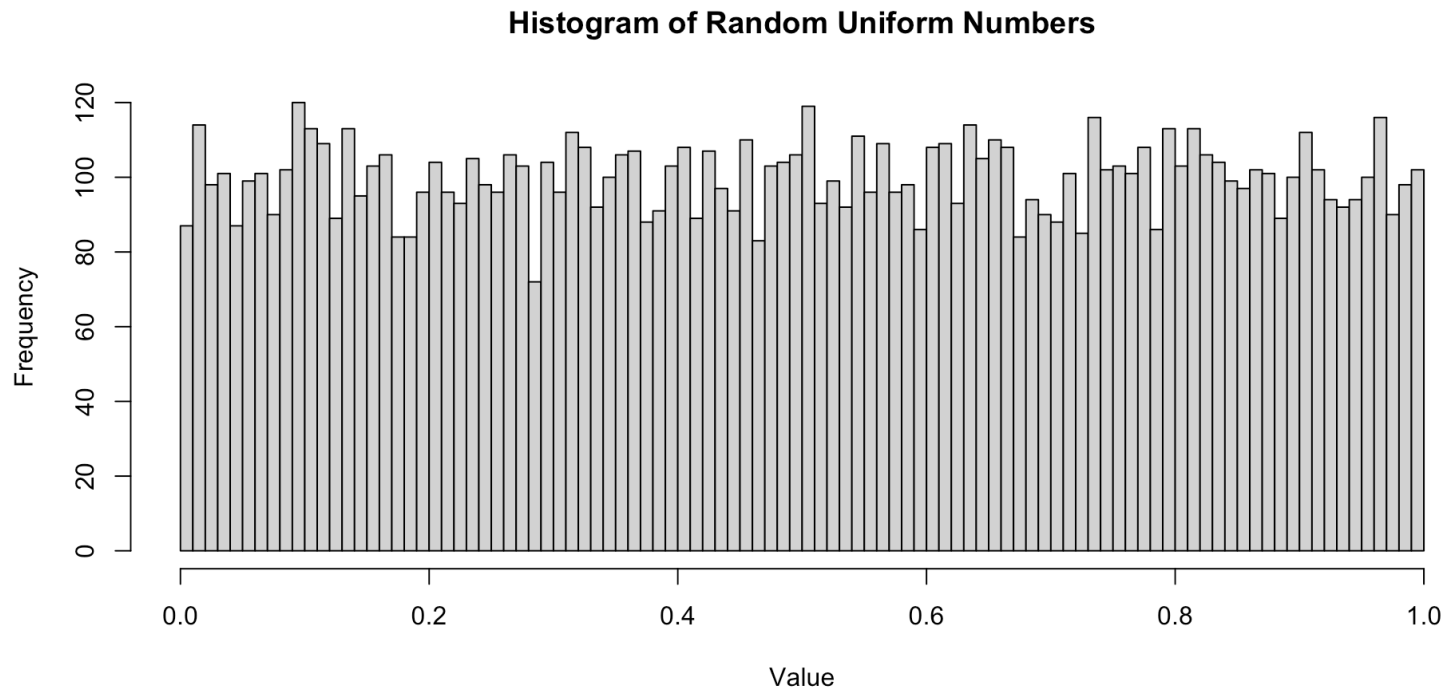
NORMAL DISTRIBUTION

```
1 # Draw 10000 random numbers from a normal distribution
2 normals <- rnorm(10000, mean = 0, sd = 1)
3 # Plot the histogram of the random numbers
4 hist(normals,
5       breaks = 30,
6       main = "Histogram of Random Normal Numbers",
7       xlab = "Value",
8       ylab = "Frequency")
```



UNIFORM DISTRIBUTION

```
1 # Draw 10000 random numbers from a uniform distribution
2 uniforms <- runif(10000, min = 0, max = 1)
3 # Plot the histogram of the random numbers
4 hist(uniforms,
5     breaks = 80,
6     main = "Histogram of Random Uniform Numbers",
7     xlab = "Value",
8     ylab = "Frequency")
```



RANDOM NUMBER GENERATORS

HOW PRNGS WORK

Pseudo Random Number Generators (PRNGs) are algorithms that generate sequences of numbers that approximate the properties of random numbers. They are called “pseudo” because they are deterministic and can be reproduced if the initial state (seed) is known.

```
1 # fix the seed
2 set.seed(123)
3 # draw 10 random integers between 1 and 100 without replacement
4 sample(1:100, size = 10, replace = FALSE)
```

```
[1] 31 79 51 14 67 42 50 43 97 25
```

```
1 # fix the seed again
2 set.seed(123)
3 # draw 10 random integers between 1 and 100 without replacement
4 sample(1:100, size = 10, replace = FALSE)
```

```
[1] 31 79 51 14 67 42 50 43 97 25
```

```
1 # draw the same 10 random numbers in sets of 5
2 set.seed(123)
3 sample(1:100, size = 5, replace = FALSE)
```

```
[1] 31 79 51 14 67
```

```
1 sample(1:100, size = 5, replace = FALSE)
```

```
[1] 42 50 43 14 25
```

BEWARE: once you fix the random seed, everything that uses random numbers will become seed-dependent. Your findings can be accidental. Always replicate with another seed!

HOW PRNGS WORK

Pseudo Random Number Generators (PRNGs) are algorithms that generate sequences of numbers that approximate the properties of random numbers. They are called “pseudo” because they are deterministic and can be reproduced if the initial state (seed) is known.

```
1 # draw 10 numbers
2 set.seed(123)
3 rnorm(10)
```

```
[1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774  1.71506499
[7]  0.46091621 -1.26506123 -0.68685285 -0.44566197
```

```
1 # draw 10 numbers in sets of 5
2 set.seed(123)
3 rnorm(5)
```

```
[1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
```

```
1 rnorm(5)
```

```
[1]  1.7150650  0.4609162 -1.2650612 -0.6868529 -0.4456620
```

```
1 # draw 15 numbers in two sets, where the first set is 5 numbers
2 set.seed(123)
3 rnorm(5)
```

```
[1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774
```

```
1 rnorm(10)
```

```
[1]  1.7150650  0.4609162 -1.2650612 -0.6868529 -0.4456620  1.2240818
[7]  0.3598138  0.4007715  0.1106827 -0.5558411
```

REPLICATION VS REPRODUCTION

Reproduction is the process of running the same analysis with the same data and code to see if the results can be exactly replicated.

```
1 # reproduction
2 set.seed(123)
3 rnorm(10) %>% mean()
```

```
[1] 0.07462564
```

```
1 set.seed(123)
2 rnorm(10) %>% mean()
```

```
[1] 0.07462564
```

Replication is the process of running the same analysis on a different dataset or in a different context to see if the results are consistent.

```
1 # replication WITH reproduction
2 set.seed(123)
3 rnorm(10) %>% mean()
```

```
[1] 0.07462564
```

```
1 set.seed(124)
2 rnorm(10) %>% mean()
```

```
[1] 0.2147669
```

PRACTICAL